

## Assignment 5

Shivani Panchiwala

UTA ID: 1001982478

```
# Professor allowed me to use other small dataset because my laptop is
getting stuck while implement Food101 dataset which large.
# Dataset - https://www.kaggle.com/datasets/prasunroy/natural-images
#https://www.kaggle.com/code/androbomb/using-cnn-to-classify-images-w-
pytorch/notebook
```

### Basic CNN Architecture

```
Net(
  (conv1): Conv2d(3, 12, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
  (conv2): Conv2d(12, 24, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  (drop): Dropout2d(p=0.2, inplace=False)
  (fc): Linear(in_features=24576, out_features=8, bias=True)
)
```

As a optimizer, I decided to use the **ADAM (ADaptive Moment estimation)** optimization algorithm, that is an extension to stochastic gradient descent that has recently seen broader adoption for deep learning applications in computer vision and natural language processing. For more on optimizers in PyTorch.

The training function we need to define needs the following steps:

1. Set the model to training mode;
2. Process the images in batches; we will iterate over images in batches. Inside each batches, we have to
  - A. Import labels and features;
  - B. Reset the optimizer
  - C. Push the data forward through the layers of the model
  - D. compute the loss
  - E. Backpropagate
3. Compute the Average Loss of the Model during the Epoch

We will thus calling each once per epoch.

I took epoch =10

Epoch: 10



Training batch 1 Loss: 0.215580  
Training batch 2 Loss: 0.400570  
Training batch 3 Loss: 0.431857  
Training batch 4 Loss: 0.457938  
Training batch 5 Loss: 0.393525  
Training batch 6 Loss: 0.371742  
Training batch 7 Loss: 0.244514  
Training batch 8 Loss: 0.454677  
Training batch 9 Loss: 0.197725  
Training batch 10 Loss: 0.324622  
Training batch 11 Loss: 0.438288  
Training batch 12 Loss: 0.288834  
Training batch 13 Loss: 0.297090  
Training batch 14 Loss: 0.244109  
Training batch 15 Loss: 0.429189  
Training batch 16 Loss: 0.376998  
Training batch 17 Loss: 0.506088  
Training batch 18 Loss: 0.253847  
Training batch 19 Loss: 0.438082  
Training batch 20 Loss: 0.878997  
Training batch 21 Loss: 0.137446  
Training batch 22 Loss: 0.544574  
Training batch 23 Loss: 0.506861  
Training batch 24 Loss: 0.316770  
Training batch 25 Loss: 0.412908  
Training batch 26 Loss: 0.513011  
Training batch 27 Loss: 0.392239  
Training batch 28 Loss: 0.492679  
Training batch 29 Loss: 0.127278  
Training batch 30 Loss: 0.724068  
Training batch 31 Loss: 0.575240  
Training batch 32 Loss: 0.350005  
Training batch 33 Loss: 0.330589  
Training batch 34 Loss: 0.306604  
Training batch 35 Loss: 0.307523  
Training batch 36 Loss: 0.918360  
Training batch 37 Loss: 0.270447  
Training batch 38 Loss: 0.090224  
Training batch 39 Loss: 0.477621  
Training batch 40 Loss: 0.277689  
Training batch 41 Loss: 0.325278  
Training batch 42 Loss: 0.275837  
Training batch 43 Loss: 0.253396  
Training batch 44 Loss: 0.292126  
Training batch 45 Loss: 0.354066  
Training batch 46 Loss: 0.300094  
Training batch 47 Loss: 0.556521  
Training batch 48 Loss: 0.332499  
Training batch 49 Loss: 0.488544  
Training batch 50 Loss: 0.533675  
Training batch 51 Loss: 0.399599  
Training batch 52 Loss: 0.426033  
Training batch 53 Loss: 0.563066

Training batch 54 Loss: 0.299225  
Training batch 55 Loss: 0.596873  
Training batch 56 Loss: 0.575532  
Training batch 57 Loss: 0.283631  
Training batch 58 Loss: 0.285560  
Training batch 59 Loss: 0.483531  
Training batch 60 Loss: 0.423469  
Training batch 61 Loss: 0.479524  
Training batch 62 Loss: 0.242024  
Training batch 63 Loss: 0.340023  
Training batch 64 Loss: 0.303990  
Training batch 65 Loss: 0.273993  
Training batch 66 Loss: 0.368834  
Training batch 67 Loss: 0.582135  
Training batch 68 Loss: 0.323422  
Training batch 69 Loss: 0.252664  
Training batch 70 Loss: 0.242119  
Training batch 71 Loss: 0.367343  
Training batch 72 Loss: 0.460398  
Training batch 73 Loss: 0.670769  
Training batch 74 Loss: 0.251815  
Training batch 75 Loss: 0.220071  
Training batch 76 Loss: 0.251227  
Training batch 77 Loss: 0.320876  
Training batch 78 Loss: 0.402557  
Training batch 79 Loss: 0.371436  
Training batch 80 Loss: 0.368226  
Training batch 81 Loss: 0.486412  
Training batch 82 Loss: 0.267487  
Training batch 83 Loss: 0.209389  
Training batch 84 Loss: 0.282732  
Training batch 85 Loss: 0.451458  
Training batch 86 Loss: 0.276844  
Training batch 87 Loss: 0.272532  
Training batch 88 Loss: 0.542886  
Training batch 89 Loss: 0.355977  
Training batch 90 Loss: 0.370555  
Training batch 91 Loss: 0.375482  
Training batch 92 Loss: 0.196394  
Training batch 93 Loss: 0.518548  
Training batch 94 Loss: 0.567814  
Training batch 95 Loss: 0.235751  
Training batch 96 Loss: 0.338469  
Training batch 97 Loss: 0.328206

Training set: Average loss: 0.380777

Validation set: Average loss: 0.331784, Accuracy: 1798/2070 (87%)

Getting predictions from test set...



All Convolutional Network

Network used:

```

Allconvos(
  (conv1): Conv2d(3, 16, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (conv4): Conv2d(64, 64, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (dropout): Dropout2d(p=0.4, inplace=False)
  (batchnorm1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (batchnorm2): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (batchnorm3): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (fc1): Linear(in_features=1600, out_features=512, bias=True)
  (fc2): Linear(in_features=512, out_features=256, bias=True)
  (fc3): Linear(in_features=256, out_features=4, bias=True)
)

```

Epoch: 10

```

Training batch 1 Loss: 0.256552
Training batch 2 Loss: 0.277665
Training batch 3 Loss: 0.371196
Training batch 4 Loss: 0.259750
Training batch 5 Loss: 0.177695
Training batch 6 Loss: 0.243364
Training batch 7 Loss: 0.221827
Training batch 8 Loss: 0.362315
Training batch 9 Loss: 0.191053
Training batch 10 Loss: 0.198967
Training batch 11 Loss: 0.244452
Training batch 12 Loss: 0.339923
Training batch 13 Loss: 0.241270
Training batch 14 Loss: 0.213669
Training batch 15 Loss: 0.314441
Training batch 16 Loss: 0.271592
Training batch 17 Loss: 0.311054
Training batch 18 Loss: 0.250080
Training batch 19 Loss: 0.337135
Training batch 20 Loss: 0.244453
Training batch 21 Loss: 0.214220
Training batch 22 Loss: 0.457922
Training batch 23 Loss: 0.407007
Training batch 24 Loss: 0.334117
Training batch 25 Loss: 0.275229
Training batch 26 Loss: 0.283150
Training batch 27 Loss: 0.222227
Training batch 28 Loss: 0.333721
Training batch 29 Loss: 0.251266
Training batch 30 Loss: 0.343995
Training batch 31 Loss: 0.240108
Training batch 32 Loss: 0.151909
Training batch 33 Loss: 0.284326

```

```
Training batch 34 Loss: 0.348609
Training batch 35 Loss: 0.295026
Training batch 36 Loss: 0.381754
Training batch 37 Loss: 0.287323
Training batch 38 Loss: 0.258291
Training batch 39 Loss: 0.287313
Training batch 40 Loss: 0.250249
Training batch 41 Loss: 0.384035
Training batch 42 Loss: 0.285144
Training batch 43 Loss: 0.173163
Training batch 44 Loss: 0.210336
Training batch 45 Loss: 0.319180
Training batch 46 Loss: 0.167396
Training batch 47 Loss: 0.346009
Training batch 48 Loss: 0.316399
Training batch 49 Loss: 0.254362
Training batch 50 Loss: 0.267434
Training batch 51 Loss: 0.313971
Training batch 52 Loss: 0.292695
Training batch 53 Loss: 0.521185
Training batch 54 Loss: 0.375270
Training batch 55 Loss: 0.377329
Training batch 56 Loss: 0.404574
Training batch 57 Loss: 0.156820
Training batch 58 Loss: 0.140445
Training batch 59 Loss: 0.371543
Training batch 60 Loss: 0.257344
Training batch 61 Loss: 0.439362
Training batch 62 Loss: 0.177520
Training batch 63 Loss: 0.338472
Training batch 64 Loss: 0.344273
Training batch 65 Loss: 0.273731
Training batch 66 Loss: 0.254769
Training batch 67 Loss: 0.194062
Training batch 68 Loss: 0.290815
Training batch 69 Loss: 0.287370
Training batch 70 Loss: 0.330671
Training batch 71 Loss: 0.414598
Training batch 72 Loss: 0.422282
Training batch 73 Loss: 0.327267
Training batch 74 Loss: 0.151157
Training batch 75 Loss: 0.225146
Training batch 76 Loss: 0.242818
Training batch 77 Loss: 0.279342
Training batch 78 Loss: 0.291669
Training batch 79 Loss: 0.196575
Training batch 80 Loss: 0.275824
Training batch 81 Loss: 0.314042
Training batch 82 Loss: 0.144787
Training batch 83 Loss: 0.327011
Training batch 84 Loss: 0.233003
Training batch 85 Loss: 0.147416
Training batch 86 Loss: 0.241256
Training batch 87 Loss: 0.250847
Training batch 88 Loss: 0.156115
Training batch 89 Loss: 0.394556
Training batch 90 Loss: 0.294581
Training batch 91 Loss: 0.292210
```

```
Training batch 92 Loss: 0.105365
Training batch 93 Loss: 0.265292
Training batch 94 Loss: 0.115586
Training batch 95 Loss: 0.178074
Training batch 96 Loss: 0.165241
Training batch 97 Loss: 0.312843
Training set: Average loss: 0.276977
Validation set: Average loss: 0.274173, Accuracy: 1837/2070 (89%)
```

In All convolutional is more accurate then Basic CNN.

## L1 Regularization

I took only 5 epoch

```
Starting epoch 5
Loss after mini-batch 500: 2.94215 (of which 0.63957 L1 loss)
Loss after mini-batch 1000: 2.93445 (of which 0.63186 L1 loss)
Loss after mini-batch 1500: 2.93992 (of which 0.63733 L1 loss)
Loss after mini-batch 2000: 2.94098 (of which 0.63839 L1 loss)
Loss after mini-batch 2500: 2.93725 (of which 0.63468 L1 loss)
Loss after mini-batch 3000: 2.93692 (of which 0.63433 L1 loss)
Loss after mini-batch 3500: 2.94259 (of which 0.63999 L1 loss)
Loss after mini-batch 4000: 2.93637 (of which 0.63378 L1 loss)
Loss after mini-batch 4500: 2.93862 (of which 0.63604 L1 loss)
Loss after mini-batch 5000: 2.93913 (of which 0.63653 L1 loss)
Loss after mini-batch 5500: 2.93972 (of which 0.63714 L1 loss)
Loss after mini-batch 6000: 2.94161 (of which 0.63903 L1 loss)
Training process has finished.
```

## L2 Regularization

```
Starting epoch 5
Loss after mini-batch 500: 2.30324 (of which 0.00011 l2 loss)
Loss after mini-batch 1000: 2.30488 (of which 0.00011 l2 loss)
Loss after mini-batch 1500: 2.30274 (of which 0.00004 l2 loss)
Loss after mini-batch 2000: 2.30191 (of which 0.00004 l2 loss)
Loss after mini-batch 2500: 2.30102 (of which 0.00007 l2 loss)
Loss after mini-batch 3000: 2.30194 (of which 0.00008 l2 loss)
Loss after mini-batch 3500: 2.30228 (of which 0.00007 l2 loss)
Loss after mini-batch 4000: 2.30383 (of which 0.00009 l2 loss)
Loss after mini-batch 4500: 2.30140 (of which 0.00006 l2 loss)
Loss after mini-batch 5000: 2.30330 (of which 0.00009 l2 loss)
Loss after mini-batch 5500: 2.30312 (of which 0.00007 l2 loss)
Loss after mini-batch 6000: 2.30295 (of which 0.00011 l2 loss)
Training process has finished.
```

Transfer Learning

```
#I did this whole assignmet using transfer learning....
```

```
!unzip archive\ \1\).zip
# The images are in a folder named 'input/natural-
images/natural_images'
training_folder_name = 'data/natural_images'

# All images are 128x128 pixels
img_size = (128,128)

# The folder contains a subfolder for each class of shape
classes = sorted(os.listdir(training_folder_name))
print(classes)
```

```
['airplane', 'car', 'cat', 'dog', 'flower', 'fruit', 'motorbike',
'person']
```

```
from PIL import Image

# function to resize image
def resize_image(src_image, size=(128,128), bg_color="white"):
    from PIL import Image, ImageOps

    # resize the image so the longest dimension matches our target size
    src_image.thumbnail(size, Image.ANTIALIAS)

    # Create a new square background image
    new_image = Image.new("RGB", size, bg_color)

    # Paste the resized image into the center of the square background
    new_image.paste(src_image, (int((size[0] - src_image.size[0]) / 2),
int((size[1] - src_image.size[1]) / 2)))

    # return the resized image
    return new_image
```

```
training_folder_name = 'data/natural_images'

# New location for the resized images
train_folder = '../working/data/natural_images'

# Create resized copies of all of the source images
size = (128,128)

# Create the output folder if it doesn't already exist
```



```

if os.path.exists(train_folder):
    shutil.rmtree(train_folder)

# Loop through each subfolder in the input folder
print('Transforming images...')
for root, folders, files in os.walk(training_folder_name):
    for sub_folder in folders:
        print('processing folder ' + sub_folder)
        # Create a matching subfolder in the output dir
        saveFolder = os.path.join(train_folder, sub_folder)
        if not os.path.exists(saveFolder):
            os.makedirs(saveFolder)
        # Loop through the files in the subfolder
        file_names = os.listdir(os.path.join(root, sub_folder))
        for file_name in file_names:
            # Open the file
            file_path = os.path.join(root, sub_folder, file_name)
            #print("reading " + file_path)
            image = Image.open(file_path)
            # Create a resized version and save it
            resized_image = resize_image(image, size)
            saveAs = os.path.join(saveFolder, file_name)
            #print("writing " + saveAs)
            resized_image.save(saveAs)

print('Done.')

```

```

Transforming images...
processing folder fruit
processing folder person
processing folder dog
processing folder airplane
processing folder car
processing folder flower
processing folder cat
processing folder motorbike
Done.

```

After that I took same cnn model as I mentioned above..